

Optimization, VAE, and GAN

Seanie Lee

MLAI, KAIST

lsnfamily@kaist.ac.kr

In this tutorial,

- We do not cover optimization – Choose whatever optimizer (SGD, Adam) you want.
- Train VAE and GAN to generate MNIST digits.
- Code is available at : https://github.com/seanie12/vae_gan


Setup

- First download the code from github repository
 - Open your terminal and go to your working directory.
 - `git clone https://github.com/seanie12/vae_gan.git`
- If you want to run the code in colab, upload the vae.ipynb or gan.ipynb
- Code: https://github.com/seanie12/vae_gan



Variational Auto-Encoder (VAE)

VAE Overview



$$\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{z})] - D_{\text{KL}} [q_{\phi}(\mathbf{z}|\mathbf{x}) || p_{\theta}(\mathbf{z})]$$


Implementation TODO

- Encoder & Decoder Network
- Loss function
- Training Loop
- Visualization

Implementation TODO


- **Encoder & Decoder Network**
- Loss function
- Training Loop
- Visualization

Encoder & Decoder




$$\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{z})] - D_{\text{KL}} [q_{\phi}(\mathbf{z}|\mathbf{x}) || p_{\theta}(\mathbf{z})]$$

Encoder & Decoder



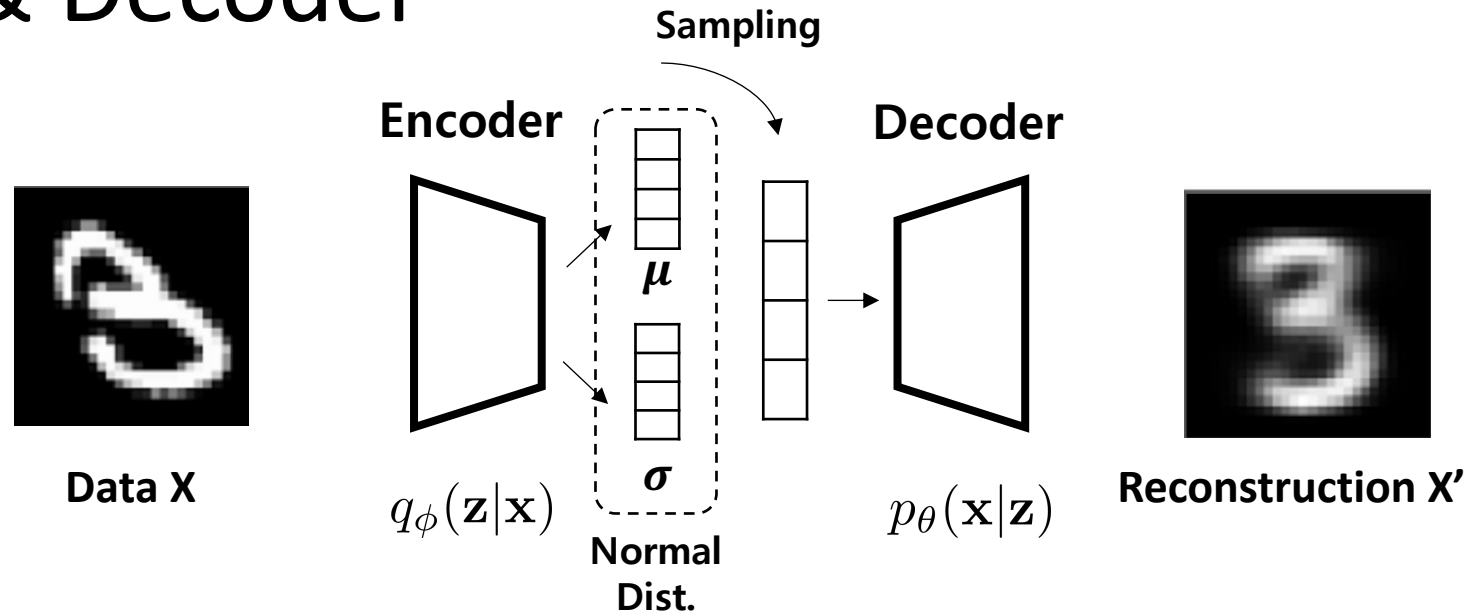
$$\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{z})] - D_{\text{KL}} [q_{\phi}(\mathbf{z}|\mathbf{x}) || p_{\theta}(\mathbf{z})]$$

Encoder & Decoder



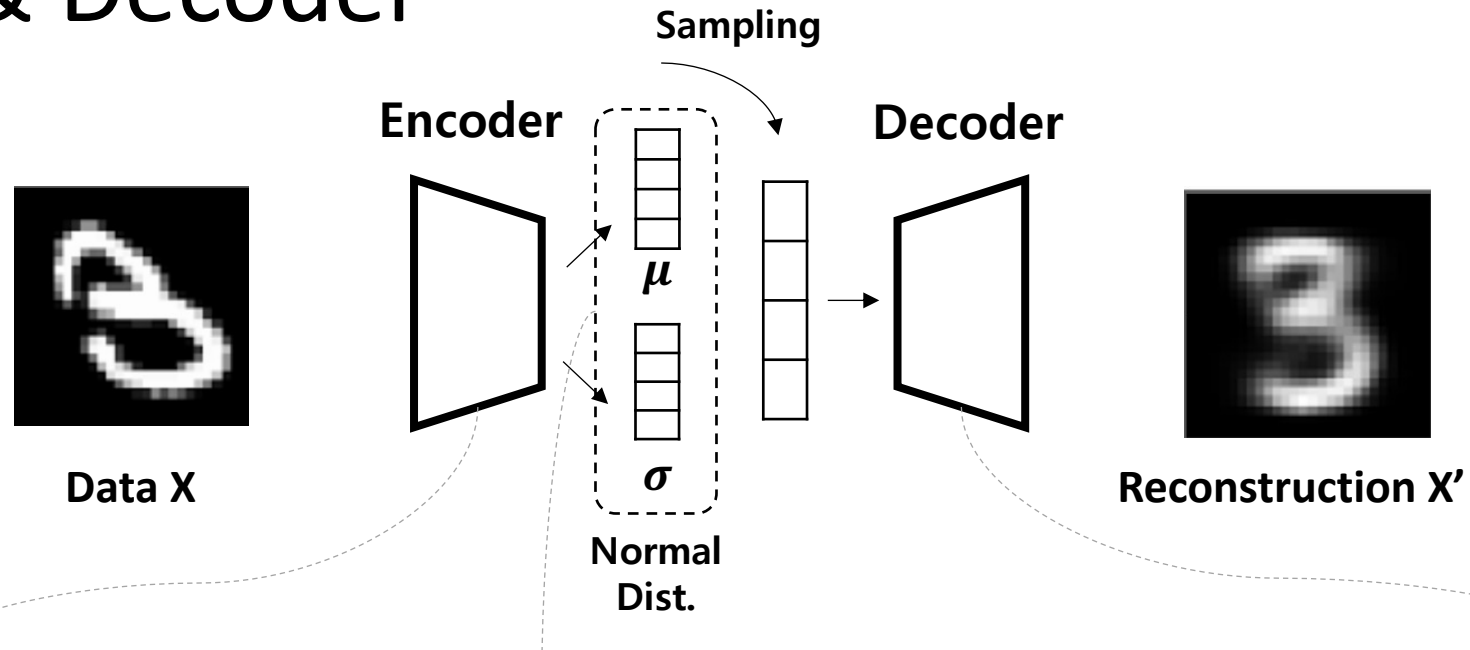
$$\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{z})] - D_{\text{KL}} [q_{\phi}(\mathbf{z}|\mathbf{x}) || p_{\theta}(\mathbf{z})]$$

Encoder & Decoder



$$\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{z})] - D_{\text{KL}} [q_{\phi}(\mathbf{z}|\mathbf{x}) || p_{\theta}(\mathbf{z})]$$

Encoder & Decoder

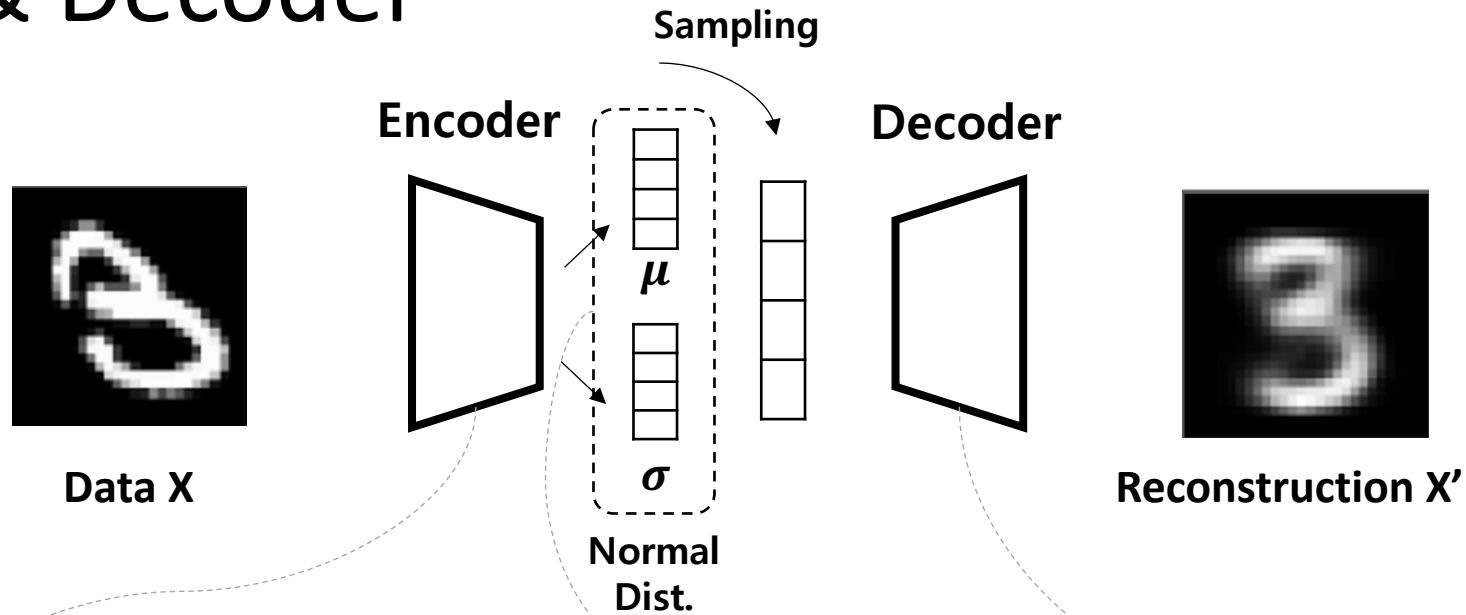


```
self.encoder = nn.Sequential(  
    nn.Linear(self.x_dim, hidden_dim),  
    nn.ReLU(),  
    nn.Linear(hidden_dim, hidden_dim),  
    nn.ReLU(),  
    nn.Linear(hidden_dim, hidden_dim),  
    nn.ReLU(),  
)
```

```
self.stat_net = nn.Linear(hidden_dim,  
                           2 * z_dim)
```

```
self.decoder = nn.Sequential(  
    nn.Linear(z_dim, hidden_dim),  
    nn.ReLU(),  
    nn.Linear(hidden_dim, hidden_dim),  
    nn.ReLU(),  
    nn.Linear(hidden_dim, self.x_dim),  
    nn.Sigmoid(),  
)
```

Encoder & Decoder



```
def forward(self, x):  
    x = x.view(-1, self.x_dim)  
    h = self.encoder(x)  
    stats = self.stat_net(h)
```

```
z_mu, z_logvar = torch.chunk(stats, 2, dim=-1)  
  
_std = (torch.randn_like(z_mu)  
        * torch.exp(0.5*z_logvar))  
z = z_mu + _std
```

```
out = self.decoder(z)
```

Implementation TODO

- Encoder & Decoder Network
- **Loss function**
- Training Loop
- Visualization

ELBO

$$\max_{\theta, \phi} \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [\log p_{\theta}(\mathbf{x}|\mathbf{z})] - D_{KL} [q_{\phi}(\mathbf{z}|\mathbf{x}) || p_{\theta}(\mathbf{z})]$$

$$\begin{aligned} \log p_{\theta}(\mathbf{x}|\mathbf{z}) &= \log \prod_{d=1}^{D_x} \text{Ber}(x_d | x'_d) \\ &= \sum_{d=1}^D \log(x'_d)^{x_d} (1 - x'_d)^{1-x_d} \\ &= \sum_{d=1}^D x_d \log x'_d + (1 - x_d) \log(1 - x'_d) \end{aligned}$$

$$\begin{aligned} D_{KL} [q_{\phi}(\mathbf{z}|\mathbf{x}) || p_{\theta}(\mathbf{z})] &= \sum_{d=1}^{D_z} D_{KL}[N(z_d | \mu_d, \sigma_d^2) || N(z_d | 0, 1)] \\ &= \frac{1}{2} \sum_{d=1}^{D_z} (\mu_d^2 + \sigma_d^2 - \log \sigma_d^2 - 1) \end{aligned}$$

```
def compute_elbo(x, x_reconst, mu, logvar):
    criterion = nn.BCELoss(reduction="sum")
    log_likelihood = -criterion(x_reconst, x.view(-1, 784)) / x.size(0)
    kl = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())

    return log_likelihood - kl
```

Implementation TODO

- Encoder & Decoder Network
- Loss function
- **Training Loop**
- Visualization

Training Loop

```
losses = []
for epoch_i in trange(n_epoch):
    for x, _ in loader:
        x = x.to(device)

        x_reconst, mu, logvar = vae(x)

        elbo = compute_elbo(x, x_reconst, mu, logvar)
        loss = -elbo


        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        losses.append(loss.item())
```



Implementation TODO

- Encoder & Decoder Network
- Loss function
- Training Loop
- **Visualization**


Visualization



```
def generate(self, z):  
    return self.decoder(z).view(len(z), *self.x_shape)
```

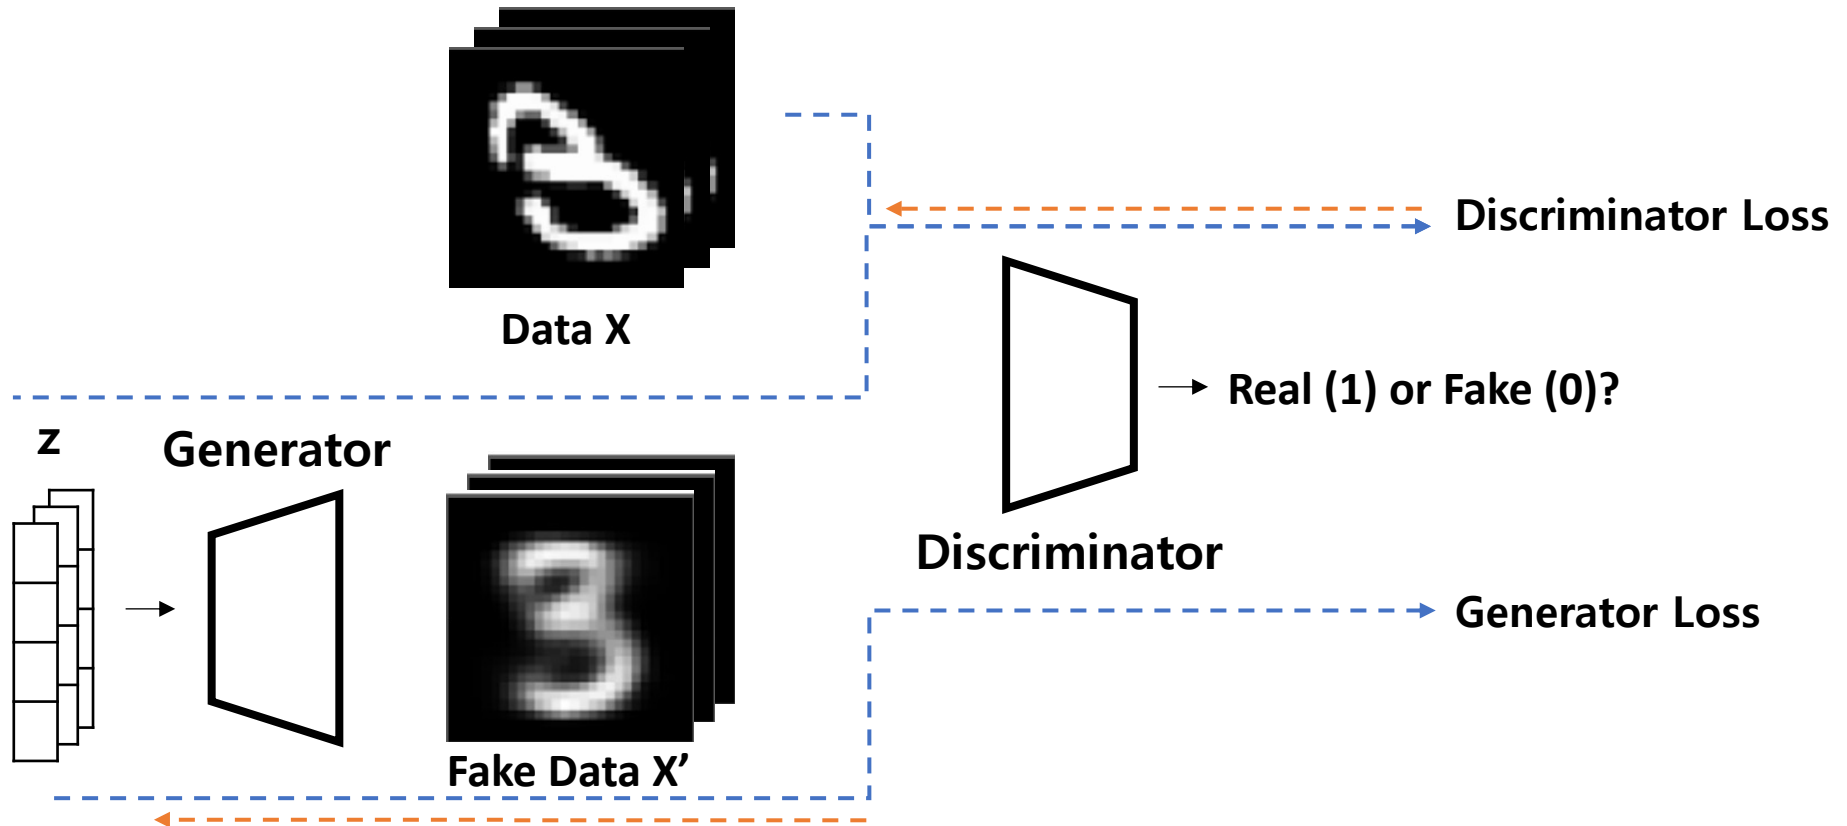
```
zs = torch.rand(16, 2).to(device)  
reconst_x = vae.generate(zs)
```

Visualization



Generative Adversarial Networks (GAN)

GAN Overview




Implementation TODO

- Generator & Discriminator Network
- Training Loop
- Visualization


Implementation TODO

- **Generator & Discriminator Network**
- Training Loop
- Visualization

Generator & Discriminator




Generator & Discriminator



```
class G(nn.Module):  
    def __init__(self, z_dim, x_shape, hidden_dim=500):  
        super().__init__()  
  
        self.x_shape = x_shape  
        self.x_dim = np.prod(x_shape)  
        self.z_dim = z_dim  
  
        self.net = nn.Sequential(  
            nn.Linear(z_dim, hidden_dim),  
            nn.ReLU(),  
            nn.Linear(hidden_dim, hidden_dim),  
            nn.ReLU(),  
            nn.Linear(hidden_dim, self.x_dim),  
            nn.Sigmoid()  
        )  
  
    def forward(self, input):  
        return self.net(input).view(len(input), *self.x_shape)
```



Generator & Discriminator



```
class D(nn.Module):
    def __init__(self, x_shape, hidden_dim=500):
        super().__init__()

        self.x_dim = np.prod(x_shape)

        self.net = nn.Sequential(
            nn.Linear(self.x_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, 1),
            nn.Sigmoid()
        )

    def forward(self, input):
        return self.net(input.view(len(input), self.x_dim)).squeeze(-1)
```

Implementation TODO

- Generator & Discriminator Network
- **Training Loop**
- Visualization

Training Loop

```
for epoch_i in trange(n_epoch):  
    for real_x, _ in loader:
```

```
        d.zero_grad()  
  
        # discriminator for real data  
        real_x = real_x.to(device)  
        real_pred_y = d(real_x)  
        real_y = make_real_y(len(real_x)).to(device)  
  
        real_d_loss = criterion(real_pred_y, real_y)  
        real_d_loss.backward()  
  
        # discriminator for fake data  
        z = torch.randn(len(real_x), z_dim, device=device)  
        fake_x = g(z)  
        fake_pred_y = d(fake_x.detach())  
        fake_y = make_fake_y(len(real_x)).to(device)  
  
        fake_d_loss = criterion(fake_pred_y, fake_y)  
        fake_d_loss.backward()  
  
        d_optimizer.step()
```


```
        # generator  
        g.zero_grad()  
        fake_pred_y = d(fake_x).view(-1)  
        g_loss = criterion(fake_pred_y, real_y)  
        g_loss.backward()  
        g_optimizer.step()
```




Implementation TODO

- Generator & Discriminator Network
- Training Loop
- **Visualization**

Epoch 0



Epoch 1



Epoch 15



Epoch 29



Q&A